Figure 5: Causal clusters, or causal topics, are shown to the right of the produced causal claim network. Each topic consists of a set of keywords that describes the cluster.

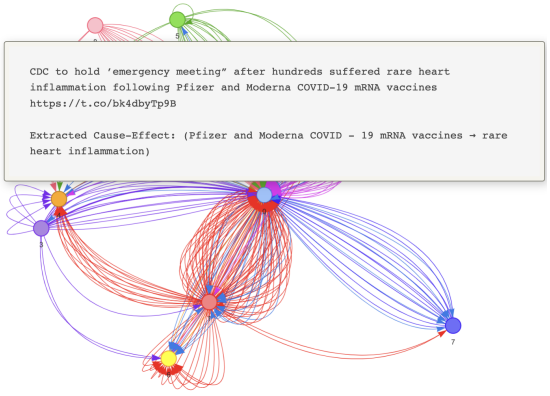

Figure 6: Hovering over an edge in the causal claim network displays the document and extracted causal claim that constitutes that edge. The document is shown at the top of the box, and the extracted cause claim is at the bottom.

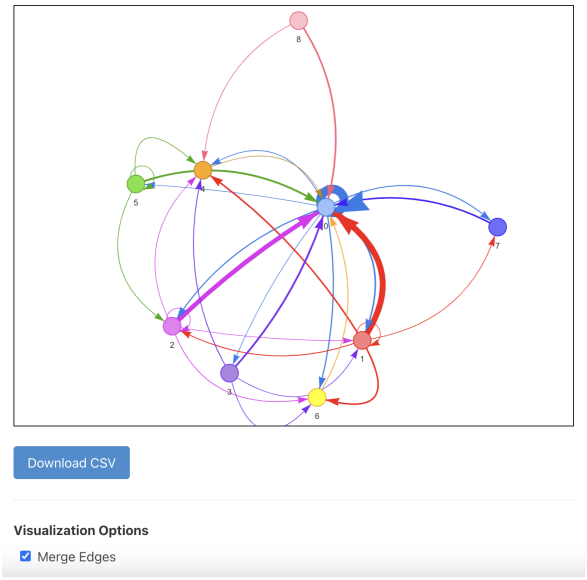

Figure 7: Causal claim network with merged edges, where edge weights equates to the number of documents linking two clusters. Merging edges is useful to quickly assess degree of linkage between causal clusters (nodes) in the network.

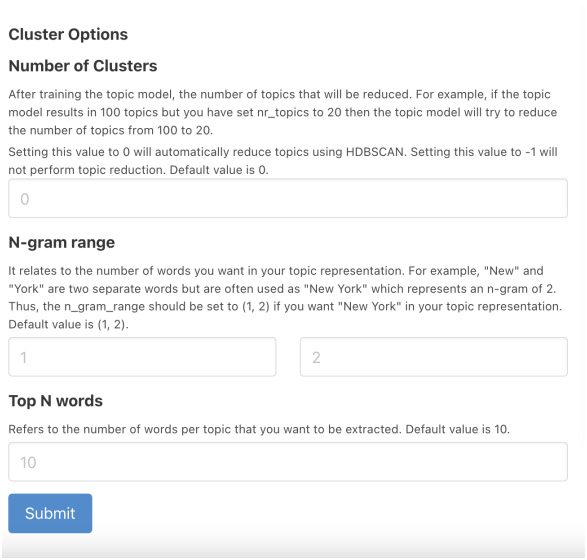

Figure 8: A user can specify parameters when running the pipeline to engage with exploratory data analysis. Users can specify the number of clusters, the n-gram range used during processing, and set the number of words to describe each topic.